

Contents

Initial Information	4
Data Dictionary	4
Initial Setup of Data	4
Including the necessary imports	4
Reading the Data in a DataFrame	5
Listing the Variables in the dataset.....	5
Solution	5
Code Result	5
Data Preprocessing Activities.....	5
Solution	6
Preprocessing Task 1 – Deal with missing values	6
Code Result	6
Observation.....	7
Pre-processing task 2 – Removing the duplicated records.....	7
Code Result	7
Observation.....	7
Code Result	7
Observation.....	7
Pre-processing task 3 – Fixing Outliers	8
Code Result	8
Code Result	9
Code Result	10
Code Result	10
Observation.....	10
Pre-processing task 4 – Datatype mismatch solution.....	10
Code Result	11
Code Result	11
Code Result	11
Observation.....	11
Code Result	11
Code Result	12

Observation.....	12
Code Result	12
Code Result	12
Observation.....	13
Visualize data to get Insights.	13
Solution	13
Insight 1.....	13
Code Plot.....	13
Observation.....	14
Code Result	14
Insight 2.....	14
Code Plot.....	15
Insight 3.....	15
Code Plot.....	15
Insight 4.....	16
Code Plot.....	16
Insight 5.....	17
Code Plot.....	18
Observation.....	18
Correlation between independent variables.....	18
Code Result	19
Code Plot.....	19
Modelling	20
Solution	20
Code Result	20
Code Result	21
Code Result	21
Observation.....	21
Code Result	22
Explaining the Approach Taken	22
State Linear Regression Equation	22
Solution	23
Key Insights from the results obtained	23

Conclusion.....	23
-----------------	----

Initial Information

The given data is the credit facility dataset that contained different features about the customers. We had to perform the different data preparation techniques to make our data ready for modelling. Once the preparation was done, the modelling was to be done using the scikit-learn linear regression model. There is a feature **B1** i.e., the customer's billable amount in 1st month which will be predicted using the linear regression model.

Data Dictionary

The list of features of the dataset is following:

ID	Customer's Unique Identifier
LIMIT	Customer's Total Limit
BALANCE	Customer's current credit balance
INCOME	Customer's current income
GENDER	Customer's Gender
EDUCATION	Customer's Highest Education attained
MARITAL	Customer's Marital Status
AGE	Customer's Age in Years
S(n)	Customer's repayment reflected status in nth month
B(n)	Customer's Billable amount in nth month
R(n)	Customer's previous repayment amount, paid in nth month
RATING	Customer's rating

Initial Setup of Data

There is some initial setup which needed to be done for the data. This setup is almost same for all the data analysis projects.

Including the necessary imports

```
# Importing required packages
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline

# Import model functions from scikit-learn
from sklearn.linear_model import LinearRegression
from sklearn.preprocessing import normalize
from sklearn.model_selection import train_test_split
```

Reading the Data in a DataFrame

```
# Reading the given data file in pandas dataframe
customers = pd.read_csv('../Data/customers.csv')
# Check the data
customers.head()
```

Listing the Variables in the dataset

At first, we done listing the categorical and numerical variables in the dataset.

Solution

There are different kinds of variables in the dataset. They can be either numeric or categorical. These two kinds are further divided into different parts. But if we talk about these types, they reflect the behavior of the variable. In our dataset, there are different categorical and the numerical variables.

From the Data Dictionary of the data provided above, we can see that there are some variables which have different classes in them. Numeric symbols are used to depict the values of those categories. So, we can call them as the categorical variables. We can display them separately. The variables which behave as the category are listed below:

```
# Defining the categorical variables
categorical_variables = ['GENDER', 'EDUCATION', 'MARITAL', 'S1', 'S2', 'S3',
                        'S4', 'S5', 'RATING']
# Defining the numerical variables
numeric_variables = [variable for variable in customers.columns if variable not
in categorical_variables]

print(f"Categorical Variables:\n{categorical_variables}\n")
print(f"Numeric Variables:\n{numeric_variables}")
```

Code Result

```
Categorical Variables:
['GENDER', 'EDUCATION', 'MARITAL', 'S1', 'S2', 'S3', 'S4', 'S5', 'RATING']

Numeric Variables:
['ID', 'LIMIT', 'BALANCE', 'INCOME', 'AGE', 'B1', 'B2', 'B3', 'B4', 'B5', 'R1', 'R2', 'R3', 'R4', 'R5']
```

In this way, the categorical and the numerical variables were listed.

Data Preprocessing Activities

Here we conducted four (4) data pre-processing tasks for the analysis of the data, explaining the results obtained.

Solution

The data must be analyzed for checking the issues in the data. These issues are needed to be resolved before the modelling process. The four data pre-processing tasks that were performed in this solution are following:

1. Deal with Missing Values
2. Removing the Duplicated Records
3. Fixing Outliers
4. Datatype Mismatch Solution

Preprocessing Task 1 – Deal with missing values

We must look for missing data fields or blank spaces (if any). This makes our data incomplete. Then, some necessary steps should be performed to deal with the missing values. The data was looked in a following way:

```
customers.isna().sum()
```

Code Result

```
ID          0
LIMIT       0
BALANCE      0
INCOME       0
RATING       0
GENDER       0
EDUCATION    13
MARITAL      38
AGE          0
S1           0
S2           0
S3           0
S4           0
S5           0
B1           0
B2           0
B3           0
B4           0
B5           0
R1           0
R2           0
R3           0
R4           0
R5           0
dtype: int64
```

We saw that we had not excessively large number of samples in our data. So instead of dropping the samples with missing values, we deal those values by filling them with the values which occurred most of the times in a variable column.

```
# Filling the blank values with the most occurred value in a column
```

```
customers['EDUCATION'].fillna(value = 2.0, inplace=True)
```

```
customers['MARITAL'].fillna(value = 2.0, inplace=True);
```

```
# Checking the updated length of concerned variables
```

```
len(customers.EDUCATION), len(customers.MARITAL)
```

As a result, the length of these columns was filled, and we had not any other missing value in our dataset remaining.

```
# Check whether there are still any missing values or not
customers.isna().sum()
```

Observation

Our data frame did not contain any missing values by this time.

Pre-processing task 2 – Removing the duplicated records

There were some of the records which had duplicated information. One of our goals while data preparation is that there should not be any duplicated data in the dataset. Appropriate code was written to check whether there were any duplicated rows in a dataset. If any, where were they located?

```
# Check the duplicate record
is_duplicated = customers.duplicated(keep='first')

# Getting indices where the duplicated records are located
is_duplicated[is_duplicated == True]
```

Code Result

```
132    True
379    True
422    True
dtype: bool
```

Observation

We can see that records on these indices contained the repeated information. We didn't want our dataset to contain the duplicated records. So, removed the duplicate samples from the data.

```
# Keeping the first occurrence of the duplicate records to have unique
information
customers.drop_duplicates(keep='first', inplace=True)
```

Let's check are there still any duplicated records?

```
is_any_duplicate = [x for x in customers.duplicated(keep='first') if x == True]
is_any_duplicate
```

Code Result

```
[]
```

Observation

Our dataframe was cleaned from the duplicated records and one more preprocessing task had done by this time.

Pre-processing task 3 – Fixing Outliers

We didn't want any unexpected values in our dataset. So, we made this sure that there were no unexpected values in any of the variables. I made a user-defined function to display all the possible values in each column.

```
# User defined function to check the unique values in each column
def check_valid_values(df):
    for column in df.columns:
        print(f"{column} Column:")
        print(f"{df[column].unique()}\n")
# Calling the function on our own data
check_valid_values(customers)
```

Code Result

```
ID Column:
[ 1 2 3 ... 18764 18765 18766]

LIMIT Column:
[210000 260000 400000 20000 180000 30000 50000 150000 60000 500000
 200000 190000 170000 280000 290000 240000 110000 140000 250000 100000
 120000 130000 10000 70000 360000 230000 340000 220000 300000 80000
 450000 40000 310000 160000 90000 330000 390000 380000 440000 410000
 320000 420000 610000 370000 480000 350000 430000 460000 490000 470000
 270000 620000 550000 510000 560000 327680 530000 660000 750000 520000
 740000 600000 700000 800000 590000 540000 570000 630000 680000 580000
 710000 640000 760000 690000 650000 780000 730000 720000 670000 160000]

BALANCE Column:
[ 0. 10928.05 65397.85 ... 33065.9 294.525 8387.575]

INCOME Column:
[235822 278481 431993 ... 30327 94607 40267]

RATING Column:
[1 0]

GENDER Column:
[1 0]

EDUCATION Column:
[1. 2. 3. 0.]

MARITAL Column:
[2. 1. 0.]

AGE Column:
[ 30 31 51 58 42 26 44 34 23 24 40 37 35 48 25 29 43 27
 38 32 54 45 49 41 47 33 50 22 53 56 28 39 63 46 52 57
 36 59 21 199 68 55 62 -1 61 64 60 80 65 66 69 67 70 79
 72 75 71 73 74]
```



```

S1 Column:
[ 0 -1  2  4  3  1  7  5  6]

S2 Column:
[ 0  2  3 -1  4  1  6  7  5  8]

S3 Column:
[ 0  2 -1  3  1  4  5  6  7]

S4 Column:
[ 0 -1  2  3  4  5  7  6]

S5 Column:
[ 0 -1  2  3  5  4  7  6]

B1 Column:
[  0  54074  343591 ...   3356  78379  48905]

B2 Column:
[  0  46407  352484 ...   2758  76304  49764]

B3 Column:
[  0  38874  338823 ...  20878  52774  36535]

B4 Column:
[  0  31324  283288 ...   5190  11855  32428]

B5 Column:
[  0  24031  185288 ...  19357  48944  15313]

R1 Column:
[  0   2000  15000 ...  2977 111784   3526]

R2 Column:
[  0  2000 14000 ...  9011  4228  8998]

R3 Column:
['0' '2000' '11500' ... '8049' '129' '1926']

R4 Column:
[  0  2000  8000 ...  8040  3319 52964]

R5 Column:
[  0  72000  7000 ... 220076 16080  7022]

```

If we look at `age` column, can age be `199` or `-1`? These values are the clear outliers of the age column which can have a great impact on our results. To make sure that how many time these values occurred, the following code was written:

```

# Check how many times the unexpected values are occurring
len(customers[customers.AGE == -1]), len(customers[customers.AGE == 199])

```

Code Result

```
(5, 5)
```

We will drop the samples with these values because they may result in false information about the data.

```

# Dropping the examples with unexpected value
customers = customers[(customers.AGE != -1) & (customers.AGE != 199)]

```

The selected samples were removed.

1st and 3rd quartiles were calculated to find the inter quartile range which helped us in analyzing the lower and upper bound of age column to detect the `outliers`.

```

first_quartile = customers.AGE.quantile(q=0.25)
third_quartile = customers.AGE.quantile(q=0.75)

```

```
iqr = third_quartile - first_quartile
```

```

lower_bound = first_quartile - 1.5 * iqr
upper_bound = third_quartile + 1.5 * iqr

print(f'First Quartile is: {first_quartile}\n'
      f'Third Quartile is: {third_quartile}\n'
      f'Interquartile range is: {iqr}\n'
      f'Lower Bound is: {lower_bound}\n'
      f'Upper Bound is: {upper_bound}')

```

Code Result

```

First Quartile is: 28.0
Third Quartile is: 41.0
Interquartile range is: 13.0
Lower Bound is: 8.5
Upper Bound is: 60.5

```

```
customers[(customers['AGE'] < lower_bound) | (customers['AGE'] > upper_bound)]
```

Code Result

	ID	LIMIT	BALANCE	INCOME	RATING	GENDER	EDUCATION	MARITAL	AGE	S1	...	B1	B2	B3	B4	B5	R1	R2	R3	R4	R5
135	135	180000	235.200	191600	1	1	3.0	1.0	63	-1	...	0	0	0	0	0	0	0	0	0	0
277	277	20000	288.750	10948	0	0	2.0	2.0	63	2	...	1650	1650	1650	1650	1650	0	0	0	0	0
575	573	180000	22388.275	210972	0	1	2.0	1.0	68	0	...	133279	127636	138164	142921	153691	12000	20000	10000	20000	5000
585	583	50000	8986.250	39470	0	0	2.0	1.0	62	-1	...	50350	49289	49058	47512	48930	1800	1800	1800	2200	1774
743	741	330000	2606.975	370467	0	0	1.0	1.0	61	0	...	19364	906	2342	1702	882	906	2342	1702	882	1940
...
17932	17930	360000	0.000	393590	1	1	3.0	1.0	73	0	...	0	0	0	0	0	0	0	0	0	0
17940	17938	510000	32737.250	537796	0	1	3.0	1.0	61	0	...	181733	192903	181801	178179	223100	17000	0	6508	50000	7000
17942	17940	360000	0.000	376811	0	1	1.0	1.0	64	0	...	4900	0	0	5640	0	0	0	5640	0	0
17944	17942	160000	13860.175	169109	0	1	3.0	1.0	74	0	...	69376	66192	16905	0	19789	2268	16905	0	19789	26442
17948	17946	500000	0.000	588381	1	1	2.0	1.0	73	0	...	2826	2652	2835	8896	3850	2652	2835	8896	3850	711

Observation

When we were observing the outliers initially, we found that we cannot have any sample with the wrong values of any variable. But when observing these samples, we can see that they do not contain inaccurate values. So, we will keep these values in our dataset. So, the records with the `age = -1` and `age = 199` are removed.

We had now acceptable age values.

Pre-processing task 4 – Datatype mismatch solution

From the data dictionary, we can see that `R(n)` is the `Customer previous repayment amount` which is paid in the nth month. We know that the amount paid should be in the `numeric` format. We have amount of 5 months. But the 3rd month has the values in the string format (object data type) and not all are numbers inside the strings. Now this issue is causing the data type mismatch issue.

Can we work with the amount in that column regardless of the data type? I tried visualize the result of this variable.

```

# try to plot the values of R3 column
try:

```

```
customers.R3.plot()
# Print the error message if it fails to plot
except TypeError:
    print("No numeric data to plot")
```

Code Result

```
No numeric data to plot
```

Oops! We got an exception error. This was because we cannot plot the data which is of non-numeric type. We also know that all the data in the dataset should be in the numeric format before any analysis. Let's identify the data type of the R3 column.

```
customers.R3.dtype
```

Code Result

```
dtype('O')
```

We see that R3 has the `'Object'` datatype. To keep record, we will keep our old unique values and the number of unique values that are present in the R3 column.

```
old_values = customers['R3'].unique()
len(customers['R3'].unique())
```

We saw what type of values are in this column apart from numbers.

```
print([x for x in customers.R3 if not str(x).isdigit()])
```

Code Result

```
['$0', '$2,620', '$6,000', '$2,200', '$390', '$2,688', '$13,069', '$7,000', '$5,000', '$2,089']
```

Observation

The above result showed us that there were not only the numbers inside the quotes. But we had to deal with the strings that contained the special symbols along with the numbers. Then I applied some regex for the conversion of data type.

```
# Replace the `$` and `,` from the Price.
customers['R3'] = customers['R3'].str.replace('[\$\,]', '', regex=True)
# Check whether there is still any unexpected value
print([x for x in customers.R3 if not str(x).isdigit()])
```

Code Result

```
[]
```

```
# keep record of new values and it's length
```

```
new_values = customers['R3'].unique()
len(customers['R3'].unique())
```

Check the difference between the old and new values

```
len(old_values) - len(new_values)
```

Code Result

9

Observation

There were 10 values in total which had different format than the numbers. But after the conversion, we are facing the difference of 9. Why is it? Is there any new value which is added to column?

```
set(new_values) - set(old_values)
```

Code Result

```
{'13069'}
```

Ah! now it makes sense. This value was only in the string format before but not in numeric. So after conversion, this was added as a numeric value.

Converting the Column into Numeric format.

```
customers['R3'] = customers['R3'].astype('int64')
# Checking the information about the data again
customers.info()
```

Code Result

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 18756 entries, 0 to 18768
Data columns (total 24 columns):
 #   Column      Non-Null Count  Dtype
---  -
 0   ID          18756 non-null  int64
 1   LIMIT       18756 non-null  int64
 2   BALANCE     18756 non-null  float64
 3   INCOME      18756 non-null  int64
 4   RATING      18756 non-null  int64
 5   GENDER      18756 non-null  int64
 6   EDUCATION   18756 non-null  float64
 7   MARITAL     18756 non-null  float64
 8   AGE         18756 non-null  int64
 9   S1          18756 non-null  int64
10  S2          18756 non-null  int64
11  S3          18756 non-null  int64
12  S4          18756 non-null  int64
13  S5          18756 non-null  int64
14  B1          18756 non-null  int64
15  B2          18756 non-null  int64
16  B3          18756 non-null  int64
17  B4          18756 non-null  int64
18  B5          18756 non-null  int64
19  R1          18756 non-null  int64
20  R2          18756 non-null  int64
21  R3          18756 non-null  int64
22  R4          18756 non-null  int64
23  R5          18756 non-null  int64
dtypes: float64(3), int64(21)
memory usage: 3.6 MB
```

Observation

Our column now contained only numeric values without any inconsistent format.

Visualize data to get Insights.

Here we articulated five (5) relevant insights of the data, with supporting visualization for each insight.

Solution

Data visualization helps us getting the useful insights from our data. We can better understand the visual results rather than struggling on the raw data. Various charts are used to help us getting familiar about the data. We performed different analysis on the data to get the proper idea about it.

Insight 1

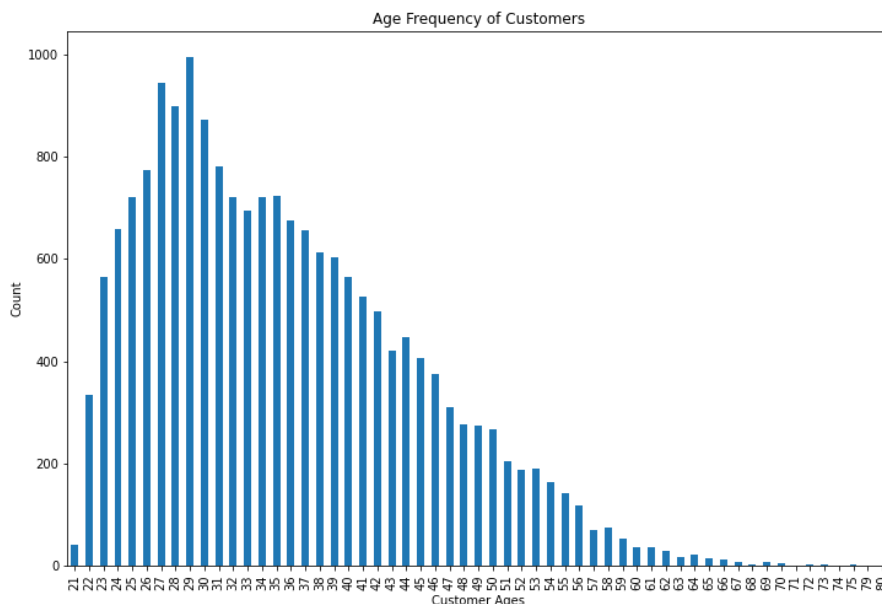
```
# Check the ages of the customers
customers['AGE'].value_counts()
```

Visualizing to see the results in better way

```
customer_ages = customers['AGE'].value_counts().sort_index()
customer_ages.plot(kind='bar', figsize=(12, 8))

plt.title("Age Frequency of Customers")
plt.xlabel("Customer Ages")
plt.ylabel("Count");
```

Code Plot



Observation

Information about the ages of the customers can be seen. We can see that there are not more old-age customers. We have more customers in the age ranging from 25-40. The interesting fact is that more customers are females. We saw that by the following command:

```
customers['GENDER'].value_counts()
```

Code Result

```
1    11598
0     7158
Name: GENDER, dtype: int64
```

We saw that we had a greater number of females in our dataset than the male customers. It means that the more samples were of female customers.

With the help of this information, we can decide what actions to take and what further facilities can be provided to the customers based on getting their age and gender dominancy.

Insight 2

If there is a need to compare two columns, we can use the function `pd.crosstab(column1, column2)`. `crosstab` is used to compute the frequency table of factors by default.

```
pd.crosstab(customers.GENDER, customers.RATING)
```

RATING	0	1
GENDER		
0	5420	1738
1	9222	2376

What can we infer from this?

Since there is the difference between the good and the bad rating ratio. We saw that there were more good ratings than the bad ratings. But the Males had almost 75-25 ratio of giving the positive and negative rating. Females have almost 80-20 ratio. It means there were more positive feedback than the negative ones.

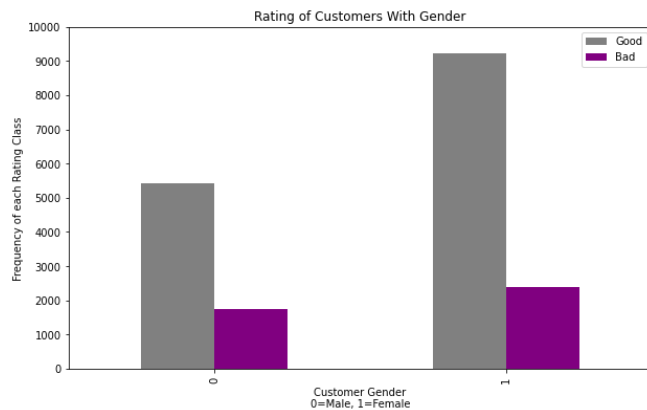
To make our crosstab visual:

```
# Plotting the crosstab
```

```
pd.crosstab(customers.GENDER, customers.RATING).plot(kind="bar",
                                                    figsize=(10,6),
                                                    color=["grey", "purple"]);
```

```
plt.legend(['Good', 'Bad']);
plt.title("Rating of Customers With Gender")
plt.xlabel("Customer Gender\n0=Male, 1=Female")
plt.ylabel("Frequency of each Rating Class")
plt.yticks(np.arange(0, 11000, 1000));
```

Code Plot



Oh, it means the customers were enjoying the facility. But the organization can also learn from this purple bar. It means that they should improve some of their terms to make this even lower. These types of visualizations would help the organization to analyze why there are negative ratings as well.

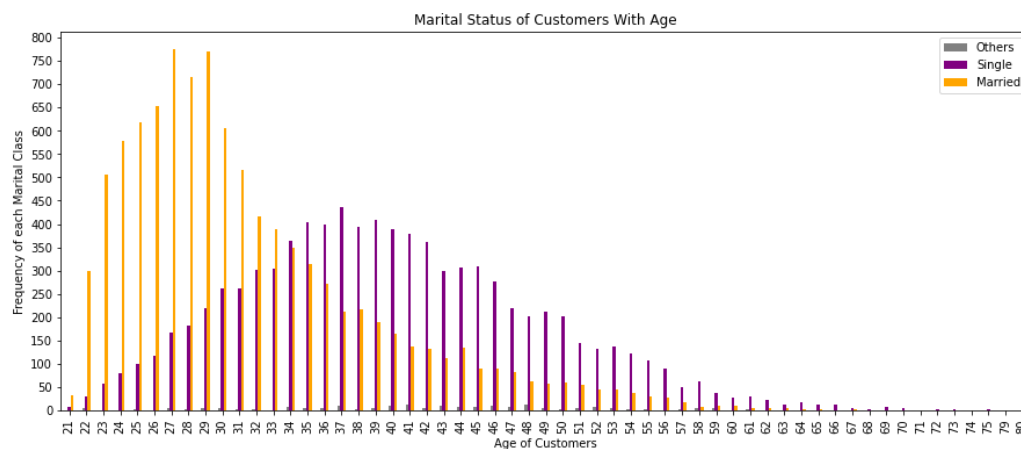
Insight 3

Now, we can do the similar analysis by identifying the marital status of the customers. In what age the customer would be married? Are there any customers who are single even in their grown age?

```
pd.crosstab(customers.AGE, customers.MARITAL).plot(kind="bar",  
                                                    figsize=(15,6),  
                                                    color=["grey", "purple",  
                                                    "orange"])
```

```
plt.legend(['Others', 'Single', 'Married']);  
plt.title("Marital Status of Customers With Age")  
plt.xlabel("Age of Customers")  
plt.ylabel("Frequency of each Marital Class")  
plt.yticks(np.arange(0, 850, 50));
```

Code Plot



Now, what we infer from here?

We can clearly see that the customers up to age 33 have the dominance of married status. But the graph has shown that the customers after age 33 are more single than the married ones. Also, there is some portion of other customers as well. But why is that?

The marital state dominance is very much based on the age based on this graph.

We can find and work based on the information we have from our visualization and do further processing if required.

Insight 4

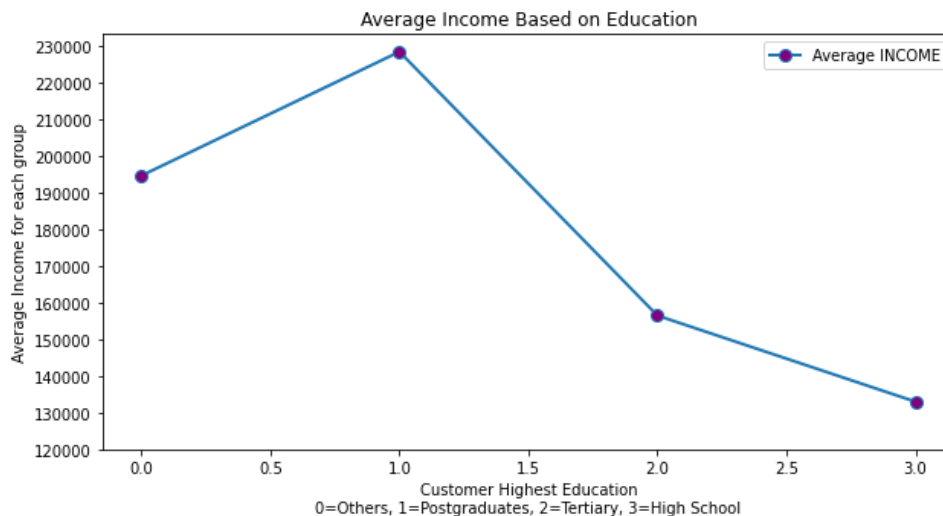
The average income of the customers based on their education was seen. Who earn more? This will also help us getting the knowledge about the income we can get based on the education level.

```
figure, axis = plt.subplots(figsize=(10, 5))

axis.plot(customers.groupby('EDUCATION')['INCOME'].mean(), marker='o', ms=8,
mfc='purple', linewidth=2, label='Average INCOME')

# Setting the title, xlabel, ylabel, values on y-axis, and legend for the plot
axis.set_title("Average Income Based on Education")
axis.set_xlabel("Customer Highest Education\n0=Others, 1=Postgraduates, 2=Tertiary, 3=High School")
axis.set_ylabel("Average Income for each group")
axis.set_yticks(np.arange(120000, 240000, 10000))
axis.legend();
```

Code Plot



Now this clearly reminds us that we are not going to get the high incomes in High School. This information clarifies the importance of education as well.

The graph told us that the highest income on average was of the customers who were 'Postgraduates'. There is also 'Other' as Education level of the customers. This plot clarifies that the other class lied between the Tertiary and the Postgraduates w.r.t. income.

Insight 5

Next, the Customer repayment reflected status of each month was found.

```
# Opening the figure with 3 rows and 2 columns. Each cell contain a subplot
fig, ((ax_0, ax_1), (ax_2, ax_3), (ax_4, ax_5)) = plt.subplots(figsize=(10, 10),
nrows=3, ncols=2)

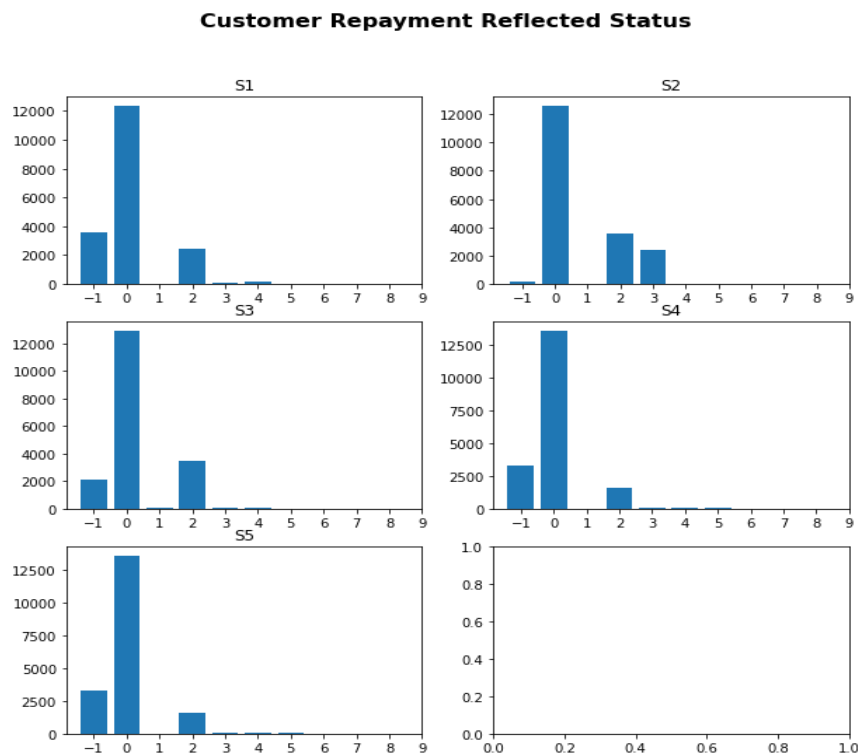
# Defining the axis to be used
axis = [ax_0, ax_1, ax_2, ax_3, ax_4]

# Using a for loop to perform the same plot on 5 different columns
for i in range(1, len(axis) + 1):
    # Making a bar plot which represent the frequency of each value
    axis[i-1].bar(customers['S'+str(i)].unique(),
customers['S'+str(i)].value_counts())

    # Adding the extra useful information in each subplot
    axis[i-1].set_title("S" + str(i))
    axis[i-1].set_xticks(np.arange(-1,10, 1))

# Setting the super title for the figure
fig.suptitle("Customer Repayment Reflected Status", fontsize=16,
fontweight='bold');
```

Code Plot



Observation

Each subplot is providing the information about the repayment reflected status of each month for the customers. The figure shows that all the customers have the 'minimum sum payment' status in maximum. It means that most of the customers has the minimum payment due specified in the statement of Accounts that should be paid by them by the Payment due date.

In most of the months, the customers with prompt payment are also greater in number.

With the help of this graph, we can get knowledge about the customers that where they are lying.

Correlation between independent variables

Finally, I compared all the independent variables in one hit.

Why?

Because this may give an idea of which independent variables may or may not have an impact on our label.

We did this using the function called `df.corr()` which created a **correlation matrix** for us, or we can say, a big table of numbers telling us how related each variable is the other.

```
correlation_matrix = customers.corr()  
correlation_matrix
```

Code Result

...

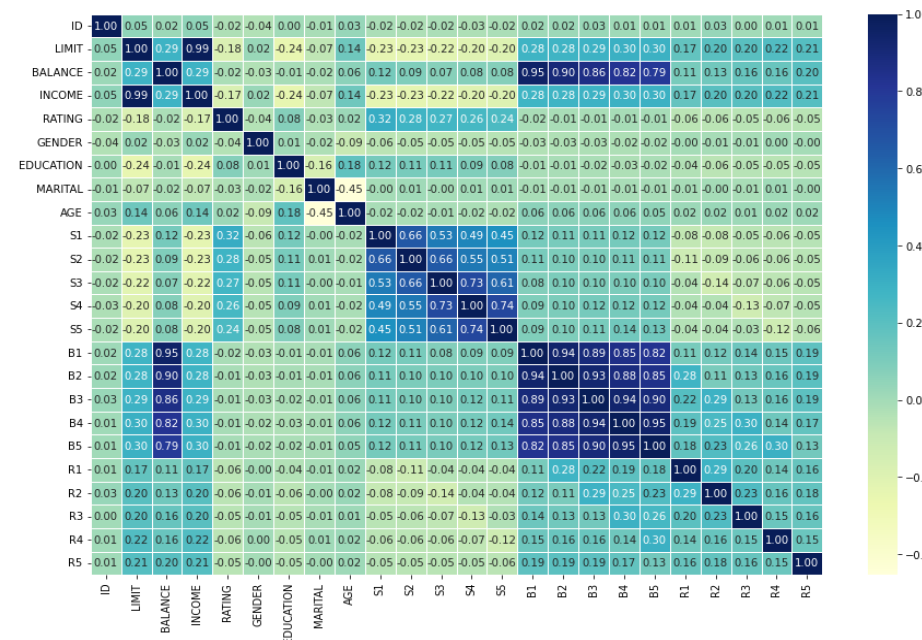
	ID	LIMIT	BALANCE	INCOME	RATING	GENDER	EDUCATION	MARITAL	AGE	S1	...	B1	B2	B3	B4	B5	R1	R2	R3	R4	R5
ID	1.000000	0.049053	0.019115	0.048510	-0.024215	-0.035078	0.002404	-0.012450	0.029862	-0.021156	...	0.016897	0.018601	0.028841	0.010694	0.010086	0.007613	0.031518	0.002365	0.007603	0.011746
LIMIT	0.049053	1.000000	0.288451	0.993748	-0.175631	0.024134	-0.243592	-0.071713	0.141387	-0.234078	...	0.279476	0.283542	0.293590	0.301587	0.297727	0.171917	0.203142	0.202466	0.222811	0.212134
BALANCE	0.019115	0.288451	1.000000	0.286335	-0.022032	-0.034711	-0.010381	-0.016615	0.057948	0.118599	...	0.950885	0.901633	0.861524	0.823847	0.794155	0.110032	0.134365	0.162067	0.156033	0.202798
INCOME	0.048510	0.993748	0.286335	1.000000	-0.173741	0.023471	-0.241721	-0.070837	0.140523	-0.231589	...	0.277372	0.281617	0.291626	0.299452	0.296145	0.169917	0.199446	0.200370	0.222412	0.210771
RATING	-0.024215	-0.175631	-0.022032	-0.173741	1.000000	-0.044543	0.078907	-0.032827	0.019423	0.323195	...	-0.015495	-0.014766	-0.012417	-0.008279	-0.006809	-0.057883	-0.059636	-0.054990	-0.061667	-0.052689
GENDER	-0.035078	0.024134	-0.034711	0.023471	-0.044543	1.000000	0.009876	-0.024352	-0.091329	-0.057033	...	-0.033734	-0.028966	-0.026764	-0.022252	-0.020417	-0.000971	-0.010547	-0.005313	0.004136	-0.002080
EDUCATION	0.002404	-0.243592	-0.010381	-0.241721	0.078907	0.009876	1.000000	-0.163941	0.178314	0.115068	...	-0.011115	-0.012805	-0.020966	-0.025325	-0.019252	-0.037796	-0.061732	-0.047183	-0.054006	-0.050473
MARITAL	-0.012450	-0.071713	-0.016615	-0.070837	-0.032827	-0.024352	-0.163941	1.000000	-0.453641	-0.001380	...	-0.011988	-0.013231	-0.012207	-0.012538	-0.008205	-0.010349	-0.003364	-0.012802	0.008214	-0.004244
AGE	0.029862	0.141387	0.057948	0.140523	0.019423	-0.091329	0.178314	-0.453641	1.000000	-0.015089	...	0.057451	0.057922	0.058300	0.055342	0.050641	0.018293	0.024190	0.013670	0.018089	0.017475
S1	-0.021156	-0.234078	0.118599	-0.231589	0.323195	-0.057033	0.115068	-0.001380	-0.015089	1.000000	...	0.116328	0.111260	0.109025	0.115306	0.118311	-0.075250	-0.084888	-0.052759	-0.058790	-0.048801
S2	-0.023941	-0.229557	0.089513	-0.228218	0.281169	-0.052007	0.112082	0.006817	-0.017270	0.664143	...	0.111645	0.102260	0.102467	0.107768	0.110735	-0.113694	-0.089547	-0.061251	-0.060768	-0.052832
S3	-0.017766	-0.219808	0.069300	-0.218470	0.268914	-0.052352	0.108911	-0.001000	-0.014722	0.532689	...	0.084994	0.099827	0.095880	0.101275	0.104121	-0.041534	-0.137692	-0.073071	-0.062283	-0.050634
S4	-0.026344	-0.202915	0.076478	-0.200902	0.257891	-0.048222	0.090867	0.005170	-0.019065	0.494103	...	0.090657	0.101831	0.119629	0.119569	0.119609	-0.035460	-0.041512	-0.125974	-0.068197	-0.046712
S5	-0.021505	-0.197564	0.075223	-0.196620	0.243755	-0.048291	0.082386	0.006819	-0.023424	0.452293	...	0.088766	0.096015	0.113171	0.136878	0.132128	-0.041235	-0.040474	-0.029514	-0.121710	-0.056361
B1	0.016897	0.279476	0.950885	0.277372	-0.015495	-0.033734	-0.011115	-0.011988	0.057451	0.116328	...	1.000000	0.939005	0.894164	0.853458	0.824070	0.113953	0.124793	0.144877	0.145382	0.187239
B2	0.018601	0.283542	0.901633	0.281617	-0.014766	-0.028966	-0.012805	-0.013231	0.057922	0.111260	...	0.939005	1.000000	0.930093	0.882172	0.848708	0.284492	0.107934	0.134621	0.155812	0.194436
B3	0.028841	0.293590	0.861524	0.291626	-0.012417	-0.026764	-0.020966	-0.012207	0.058300	0.109025	...	0.894164	0.930093	1.000000	0.935690	0.896736	0.215477	0.293126	0.130852	0.156631	0.192041
B4	0.010694	0.301587	0.823847	0.299452	-0.008279	-0.022252	-0.025325	-0.012538	0.055342	0.115306	...	0.853458	0.882172	0.935690	1.000000	0.945010	0.193043	0.247081	0.296914	0.138950	0.174449
B5	0.010086	0.297727	0.794155	0.296145	-0.006809	-0.020417	-0.019252	-0.008205	0.050641	0.118311	...	0.824070	0.848708	0.896736	0.945010	1.000000	0.179671	0.225103	0.258631	0.304237	0.126834
R1	0.007613	0.171917	0.110032	0.169917	-0.057883	-0.009971	-0.037796	-0.010349	0.018293	-0.075250	...	0.113953	0.284492	0.215477	0.193043	0.179671	1.000000	0.290046	0.198066	0.141830	0.160192
R2	0.031518	0.203142	0.134365	0.199446	-0.059636	-0.010547	-0.061732	-0.003364	0.024190	-0.084888	...	0.124793	0.107934	0.293126	0.247081	0.225103	0.290046	1.000000	0.229128	0.163520	0.176206
R3	0.002365	0.202466	0.162067	0.200370	-0.054990	-0.005313	-0.047183	-0.012802	0.013670	-0.052759	...	0.144877	0.134621	0.130852	0.296914	0.258631	0.198066	0.229128	1.000000	0.151691	0.163778
R4	0.007603	0.222811	0.156033	0.222412	-0.061667	0.004136	-0.054006	0.008214	0.018089	-0.058790	...	0.145382	0.155812	0.156631	0.138950	0.304237	0.141830	0.163520	0.151691	1.000000	0.150216
R5	0.011746	0.212134	0.202798	0.210771	-0.052689	-0.002080	-0.050473	-0.004244	0.017475	-0.048801	...	0.187239	0.194436	0.192041	0.174449	0.126834	0.160192	0.176206	0.163778	0.150216	1.000000

24 rows x 24 columns

Let's make it look a little prettier

```
correlation_matrix = customers.corr()
plt.figure(figsize=(16, 11))
sns.heatmap(correlation_matrix,
            annot=True,
            linewidths=0.5,
            fmt= ".2f",
            cmap="YlGnBu");
```

Code Plot



The visual representation of heatmap is clearer. Here, a higher positive value shows a positive correlation (direct) and a higher negative value shows a negative correlation (opposite).

Modelling

Performed linear regression modelling to predict the variable, B1, explaining the approach taken, including any further data pre-processing.

Solution

Now, Linear regression modelling was to be done on the data to predict the B1 variable. But there was a need of some further preprocessing of data before the modelling process. The preprocessing steps are done to help the model in predicting the target variable easily.

We had some variables in the data which were needed to be transformed for the better performance of the model. **Normalization** is used when there are differently scaled features. The updated values are then scaled between the range of 0 and 1.

```
# Making a list of variables which have wide range of values
variable_list = ['LIMIT', 'BALANCE', 'INCOME', 'AGE', 'B1', 'B2', 'B3', 'B4',
'B5', 'R1', 'R2', 'R3', 'R4', 'R5']

# Extracting those variables from original dataframe to separate them
variables_to_normalize = customers[variable_list]
# Keep track of column names (can be used in future)
normalized_columns = list(variables_to_normalize.columns.values)
# Keep track of row numbers (can be used in future)
normalized_rows = list(variables_to_normalize.index)
# See the extracted columns
variables_to_normalize
```

Code Result

	LIMIT	BALANCE	INCOME	AGE	B1	B2	B3	B4	B5	R1	R2	R3	R4	R5
0	210000	0.000	235822	30	0	0	0	0	0	0	0	0	0	0
1	260000	10928.050	278481	31	54074	46407	38874	31324	24031	2000	2000	2000	2000	72000
2	400000	65397.850	431993	51	343591	352484	338823	283288	185288	15000	14000	11500	8000	7000
3	20000	3695.300	22368	58	21470	200	0	28740	1295	0	0	28740	1565	0
4	180000	68.250	166900	42	1260	598	2056	4300	1802	598	2056	4300	1802	3691
...
18764	220000	33065.900	225862	39	192815	208365	88004	31237	15980	20000	5003	3047	5000	1000
18765	150000	294.525	149966	43	1828	3502	8979	5190	0	3526	8998	129	0	0
18766	30000	623.875	30327	37	3356	2758	20878	20582	19357	0	22000	4200	2000	3100
18767	80000	0.000	94607	41	78379	76304	52774	11855	48944	3409	1178	1926	52964	1804
18768	50000	8387.575	40267	46	48905	49764	36535	32428	15313	1800	1430	1000	1000	1000

18756 rows × 14 columns

Normalized the variables we selected from our dataset

```
# Perform normalization using sklearn.preprocessing.normalize on the selected
columns
normalized_array = normalize(variables_to_normalize)
```

```
# The array returned after the normalization process can be stored as the
dataframe with the rows and columns stored above
scaled_customers = pd.DataFrame(normalized_array, columns=normalized_columns,
index = normalized_rows)
# Let's see the scaled customers
scaled_customers
```

Code Result

```
...
LIMIT  BALANCE  INCOME  AGE  B1  B2  B3  B4  B5  R1  R2  R3  R4  R5
0  0.665037  0.000000  0.746811  0.000095  0.000000  0.000000  0.000000  0.000000  0.000000  0.000000  0.000000  0.000000  0.000000  0.000000
1  0.652824  0.027439  0.699228  0.000078  0.135772  0.116522  0.097607  0.078650  0.060339  0.005022  0.005022  0.005022  0.005022  0.180782
2  0.440855  0.072077  0.476115  0.000056  0.378684  0.388485  0.373429  0.312222  0.204213  0.016532  0.015430  0.012675  0.008817  0.007715
3  0.363271  0.067120  0.406282  0.001053  0.389971  0.003633  0.000000  0.522020  0.023522  0.000000  0.000000  0.522020  0.028426  0.000000
4  0.732874  0.000278  0.679537  0.000171  0.005130  0.002435  0.008371  0.017508  0.007337  0.002435  0.008371  0.017508  0.007337  0.015028
...
...
18764  0.504003  0.075751  0.517433  0.000089  0.441724  0.477348  0.201610  0.071562  0.036609  0.045818  0.011461  0.006980  0.011455  0.002291
18765  0.705490  0.001385  0.705330  0.000202  0.008598  0.016471  0.042231  0.024410  0.000000  0.016584  0.042320  0.000607  0.000000  0.000000
18766  0.500805  0.010415  0.506263  0.000618  0.056023  0.046041  0.348527  0.343585  0.323136  0.000000  0.367257  0.070113  0.033387  0.051750
18767  0.424821  0.000000  0.502388  0.000218  0.416213  0.405194  0.280244  0.062953  0.259906  0.018103  0.006255  0.010228  0.281253  0.009580
18768  0.462466  0.077579  0.372442  0.000425  0.452338  0.460283  0.337924  0.299937  0.141635  0.016649  0.013227  0.009249  0.009249  0.009249

18765 rows x 14 columns
```

We kept the original (without transformation) and the normalized variables in our dataframe to predict the target variable `B1`. Concatenation of the scaled and the original customers dataset, the columns would be merged. So, we changed the name of the transformed variables to identify them.

```
# Adding a prefix to mark the change in columns
scaled_customers = scaled_customers.add_prefix('normalized_')
# Concatenating the original and the scaled customers to predict the target
variable
updated_customers = pd.concat([customers, scaled_customers], axis=1)
# Viewing the updated customers
updated_customers
```

Code Result

```
ID  LIMIT  BALANCE  INCOME  RATING  GENDER  EDUCATION  MARITAL  AGE  S1  ...  normalized_B1  normalized_B2  normalized_B3  normalized_B4  normalized_B5  normalized_R1  normalized_R2  normalized_R3
0   1  210000   0.000  235822   1   1   1.0  2.0  30  0  ...  0.000000  0.000000  0.000000  0.000000  0.000000  0.000000  0.000000  0.000000
1   2  260000  10928.050  278481   0   0   2.0  2.0  31  0  ...  0.135772  0.116522  0.097607  0.078650  0.060339  0.005022  0.005022  0.005022
2   3  400000  65397.850  431993   0   0   3.0  1.0  51  0  ...  0.378684  0.388485  0.373429  0.312222  0.204213  0.016532  0.015430  0.012675
3   4  200000  3695.300  22368   0   0   2.0  1.0  58  -1  ...  0.389971  0.003633  0.000000  0.522020  0.023522  0.000000  0.000000  0.522020
4   5  180000   68.250  166900   0   1   2.0  1.0  42  0  ...  0.005130  0.002435  0.008371  0.017508  0.007337  0.002435  0.008371  0.017508
...  ...  ...  ...  ...  ...  ...  ...  ...  ...  ...  ...  ...  ...  ...  ...  ...  ...  ...
18764  18762  220000  33065.900  225862   0   0   3.0  1.0  39  0  ...  0.441724  0.477348  0.201610  0.071562  0.036609  0.045818  0.011461  0.006980
18765  18763  150000   294.525  149966   0   0   3.0  2.0  43  -1  ...  0.008598  0.016471  0.042231  0.024410  0.000000  0.016584  0.042320  0.000607
18766  18764  30000   623.875   30327   1   0   2.0  2.0  37  3  ...  0.056023  0.046041  0.348527  0.343585  0.323136  0.000000  0.367257  0.070113
18767  18765  80000   0.000   94607   1   0   3.0  1.0  41  -1  ...  0.416213  0.405194  0.280244  0.062953  0.259906  0.018103  0.006255  0.010228
18768  18766  50000   8387.575   40267   1   0   2.0  1.0  46  0  ...  0.452338  0.460283  0.337924  0.299937  0.141635  0.016649  0.013227  0.009249

18766 rows x 38 columns
```

Observation

Now we had a dataframe populated with the transformed and the original values of the columns. This preprocessing helped the model to predict the values in a better way then predicting the target variable with the large, ranged values.

After this, the modelling process was done on the updated data to predict the B1 variable. Appropriate code was written:

```
# Separating the features and the label in X and y
X = updated_customers.drop('B1', axis=1)
y = updated_customers['B1']

# Using the train_test_split function to split data with specifying 30% size of
test data
# Random state is specified for the reproducible results
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3,
random_state=6)

# Instantiating a model
model = LinearRegression()

# Training a model (on training portion)
model.fit(X_train, y_train)

# Make the predictions
y_preds = model.predict(X_test)
```

I compared the predicted and the true values of the data by comparing the output of y_preds and y_test. Since the model was built to predict a quantity, there was a slight difference between the model prediction and the actual values. But the model scored a good figure of **97.74%**.

```
print(f"The accuracy of the model is: {model.score(X_test, y_test)*100:.2f}%")
```

Code Result

```
The accuracy of the model is: 97.74%
```

Explaining the Approach Taken

The normalization is done on the variables with the large range of values. These type of variables causes difficulty in model to predict the target variable. When the preprocessing step is done on these variables, they are scaled into the shorter range. Hence, it becomes easy for a model to predict the target variables. So, we see the good efficiency of the model with this type of approach.

If we predict the model without scaling the large range variables, it will lag its performance. The demo was done using the code in jupyter notebook.

State Linear Regression Equation

Stated the linear regression equation and explain key insights from the results obtained.

Solution

In **linear regression**, we have X and Y (**dependent** and **independent** variables respectively). By using linear regression, we determine the behavior of dependent variable based on independent variable. There could be the positive, negative, or variable relationship between these variables.

Usually, there are many observations in the problem to be examined. Linear regression defines the line that best describes the observations. We can call that line, a regression line.

The **linear regression equation** is like the equation of line in the algebra. The equation of line is:

$$y = mx + c$$

Similarly, in linear regression we have equation:

$$\hat{y} = a_0 + a_1x$$

Where, a_0 is the y-intercept and a_1 is the slope.

y-intercept is the point where a function line passes the y-axis. We can find it by using its formula.

We usually refer **slope** as **rise over run** formula. We can find the slope and the y-intercept to determine the dependent variable for linear regression equation.

Key Insights from the results obtained

There are different observations that we can see from the result obtained from our modelling:

- Linear Regression model fits and evaluates the model to predict the quantity.
- There is a need of data transformation to scale the values with the large range of data.
- It is easy for the model to find patterns in the data with the small range, and normalization, and standardization helps in scaling the data in efficient way.
- Better results were obtained when there were variables in the dataset which had not large range values.
- There was 30% of the shuffled data on which the model was tested, and the linear model well suited on the current problem.
- When the linear regression model predicted the results, it predicted them by the linear approximation
- The linear regression model achieved more than 95% of accuracy while predicting the target variable.

Conclusion

By the process of data analysis, we can analyze the data to find the key insights from it. Based on the information we take from our analysis; we can do different decision making. There are some quality issues that may have in a data because the data is gathered from different sources by using different means to collect them. Once the data is prepared, we can perform modelling on them using some machine learning algorithm by using which the machine learning model finds the pattern in data and then predict the results to generalize its knowledge. Python provides good libraries to handle the data and create and evaluate the models.